



Queen Mary
University of London

Science and Engineering

School of Electronic Engineering and Computer Science
QMUL-BUPT Joint Programme

EBU6475 Microprocessor System Design

EBU5476 Microprocessors for Embedded Computing

Exceptions and Interrupts

References:

Chapter 4.5, 7 & 8, The Definitive Guide to ARM®;

Chapter 4 (pp 145-165; 21 pages), Embedded Systems Fundamentals

arm

Last updated: 17 March 2022

University Program Education Kits

Overview

- Exception and Interrupt Concepts
 - Entering an Exception Handler
 - Exiting an Exception Handler
- Core Interrupts
 - Using Port Module and External Interrupts
- Timing Analysis
- Program Design with Interrupts
 - Sharing Data Safely Between ISRs and Other Threads

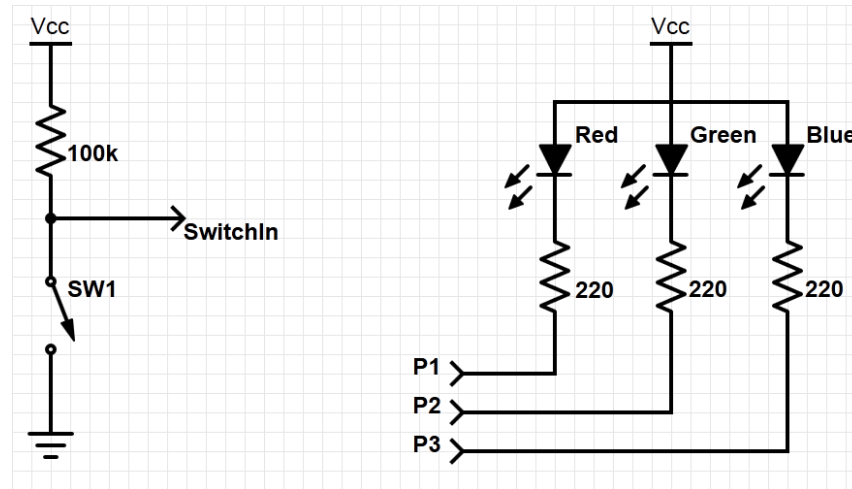
Interrupts and exceptions both alter the program flow. The difference between the two is that interrupts are used to handle external events (serial ports, keyboard) and exceptions are used to handle instruction faults, (division by zero, undefined opcode).

Interrupts are handled by the processor after finishing the current instruction. If it finds a signal on its interrupt pin, it will look up the address of the interrupt handler in the interrupt table and pass that routine control. After returning from the interrupt handler routine, it will resume program execution at the instruction after the interrupted instruction.

Exception and Interrupt Concepts

Exceptions on the other hand are divided into three kinds. These are Faults, Traps and Aborts. Faults are detected and serviced by the processor before the faulting instructions. Traps are serviced after the instruction causing the trap. User defined interrupts go into this category and can be said to be traps; this includes the MS- DOS INT 21h software interrupt, for example. Aborts are used only to signal severe system problems, when operation is no longer possible.

Example System with Interrupt




- Goal: Change color of RGB LED when switch is pressed
- Will explain details of interfacing with switch and LEDs in GPIO module later
- Need to add external switch

How to Detect Switch is Pressed?

- Polling - use software to check it
 - Slow - need to explicitly check to see if switch is pressed
 - Wasteful of CPU time - the faster a response we need, the more often we need to check
 - Scales badly - difficult to build system with many activities which can respond quickly. Response time depends on all other processing.
- Interrupt - use special hardware in MCU to detect event, run specific code (interrupt service routine - ISR) in response
 - Efficient - code runs only when necessary
 - Fast - hardware mechanism
 - Scales well
 - ISR response time doesn't depend on most other processing.
 - Code modules can be developed independently

Interrupt or Exception Processing Sequence



- Other code (background) is running
- Interrupt trigger occurs
- Processor does some hard-wired processing
- Processor executes ISR (foreground), including return-from-interrupt instruction at end
- Processor resumes other code


An **ISR** has the following key properties:

1. An interrupt request (IRQ) signal that triggers the ISR.
2. A priority level associated with the IRQ.
3. An interrupt handler function that is invoked to handle the interrupt.
4. An argument value that is passed to that function.

Interrupts

- Hardware-triggered asynchronous software routine
 - Triggered by hardware signal from peripheral or external device
 - Asynchronous - can happen anywhere in the program (unless interrupt is disabled)
 - Software routine - **Interrupt service routine** runs in response to interrupt
- Fundamental mechanism of microcontrollers
 - Provides efficient event-based processing rather than polling
 - Provides quick response to events regardless* of program state, complexity, location
 - Allows many multithreaded embedded systems to be responsive without an operating system (specifically task scheduler)

Example Program Requirements & Design



1. When Switch SW is pressed, ISR will increment count variable.
2. Main code will light LEDs according to count value in binary sequence (Blue: 4, Green: 2, Red: 1).
3. Main code will toggle its debug line each time it executes.
4. ISR will raise its debug line (and lower main's debug line) whenever it is executing.

Entering an Exception Handler

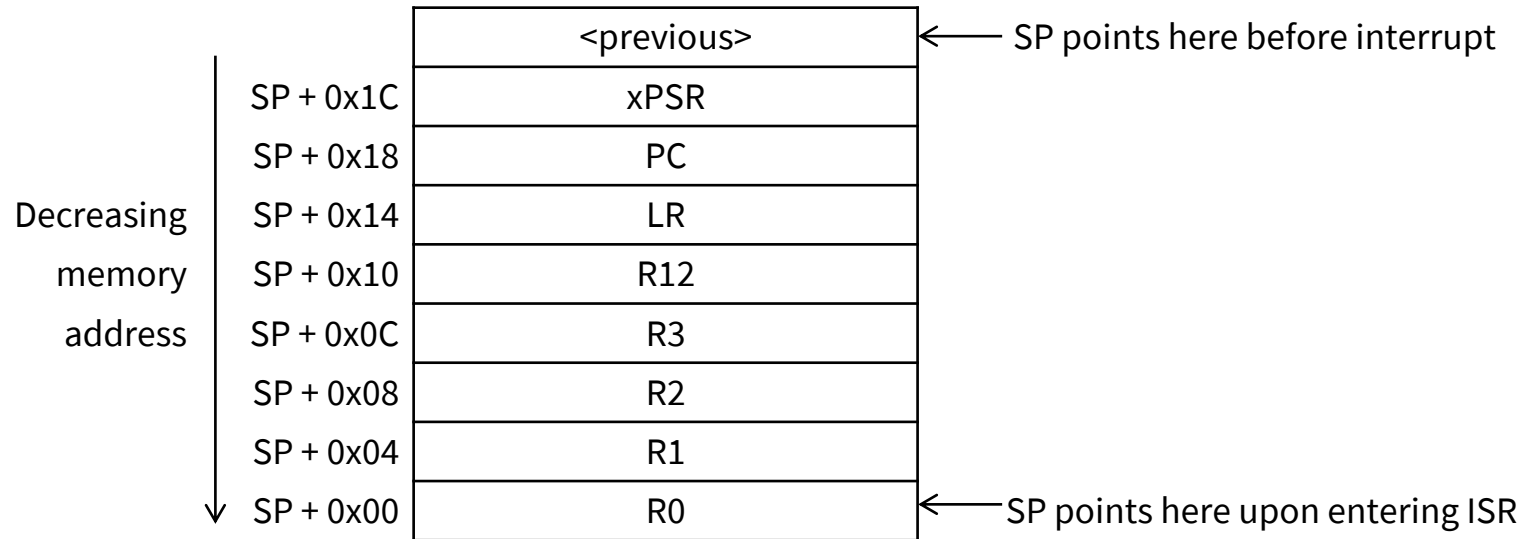
CPU's Hardwired Exception Processing

1. Finish current instruction (except for lengthy instructions)
2. Push context (8 32-bit words) onto current stack (MSP or PSP)
 - xPSR, Return address, LR (R14), R12, R3, R2, R1, R0
3. Switch to handler/privileged mode, use MSP
4. Load PC with address of exception handler
5. Load LR with EXC_RETURN code
6. Load IPSR with exception number
7. Start executing code of exception handler
8. Usually 16 cycles from exception request to execution of first instruction in handler

1. Finish Current Instruction

- Most instructions are short and finish quickly
- Some instructions may take many cycles to execute
 - Load Multiple (LDM), Store Multiple (STM), Push, Pop, MULS (32 cycles for some CPU core implementations)
- This will delay interrupt response significantly
- If one of these is executing when the interrupt is requested, the processor:
 - abandons the instruction
 - responds to the interrupt
 - executes the ISR
 - returns from interrupt
 - restarts the abandoned instruction

2. Push Context onto Current Stack



- Two SPs: Main (MSP), process (PSP)
- Which is active depends on operating mode, CONTROL register bit 1
- Stack grows toward smaller addresses.

Context Saved on Stack: uVision

| Register | Value |
|----------|------------|
| Core | |
| R0 | 0x00000000 |
| R1 | 0x00000000 |
| R2 | 0x00000000 |
| R3 | 0x0000008C |
| R4 | 0x40040000 |
| R5 | 0x00000ECC |
| R6 | 0x2EE00000 |
| R7 | 0x00000EAB |
| R8 | 0x9C7FB76F |
| R9 | 0x20000768 |
| R10 | 0x00000ECC |
| R11 | 0x00000ECC |
| R12 | 0x0FFFF0AA |
| R13 (SP) | 0x20000FD8 |
| R14 (LR) | |
| R15 (PC) | |
| xPSR | |
| Banked | |
| MSP | 0x20000FD8 |
| PSP | 0xBEF5DFF0 |
| System | |
| PRIMASK | 0 |
| CONTROL | 0x00 |
| Internal | |
| Mode | |
| Stack | |

SP value is reduced since registers have been pushed onto stack

| Address | sp |
|-------------|--|
| 0x20000FD8: | 00000000 00000000 00000000 0000008C 0FFFF0AA |
| 0x20000FEC: | 00000A0B 00000A70 21000000 00000002 00000A0B |
| 0x20001000: | 622361E2 B510BD30 6800482F 04892101 492D4308 |
| 0x20001014: | 20186008 FFD477FF 49292019 200562C8 F8EEF000 |
| 0x20001028: | 30204824 49257940 20056288 F929F000 68004823 |

Saved R0

Saved R1

Saved R2

Saved R3

Saved R12

Saved LR

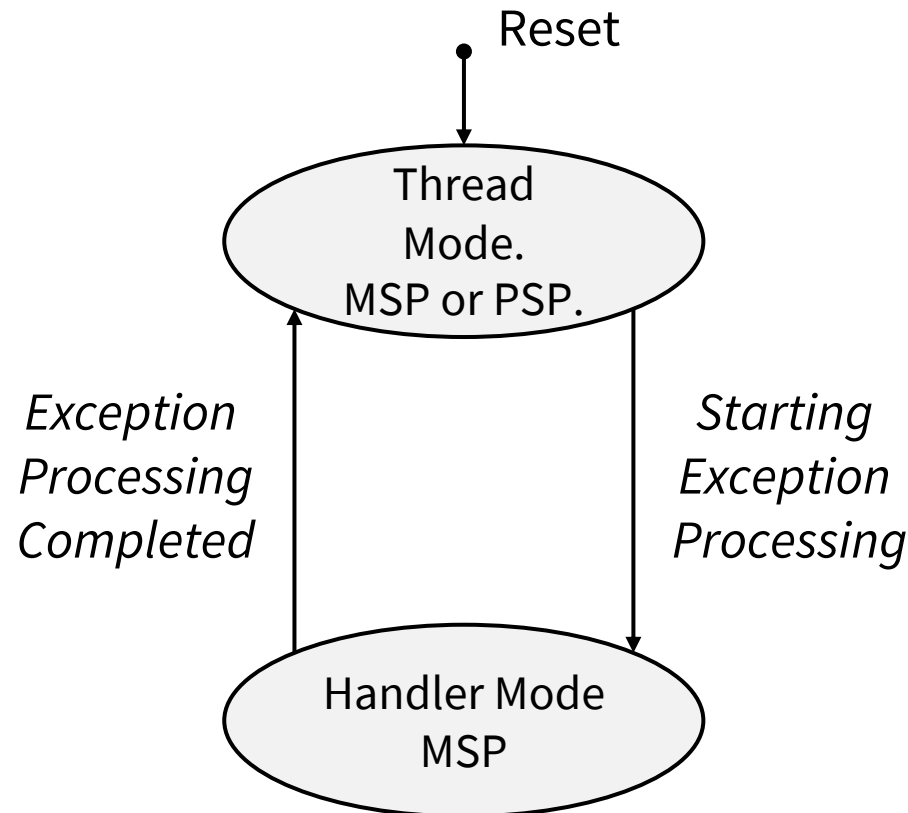
Saved PC

Saved xPSR

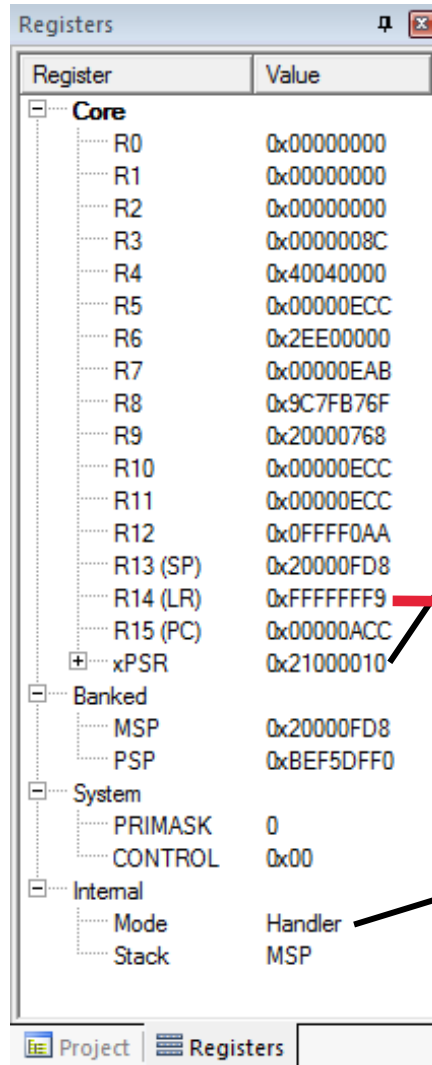
3. Switch to Handler/Privileged Mode

- Handler mode always uses Main SP

thread-normal
handler-interrupt



Handler and Privileged Mode: uVision



| Register | Value |
|-----------------|------------|
| Core | |
| R0 | 0x00000000 |
| R1 | 0x00000000 |
| R2 | 0x00000000 |
| R3 | 0x0000008C |
| R4 | 0x40040000 |
| R5 | 0x0000ECC |
| R6 | 0x2EE00000 |
| R7 | 0x0000EAB |
| R8 | 0x9C7FB76F |
| R9 | 0x20000768 |
| R10 | 0x0000ECC |
| R11 | 0x0000ECC |
| R12 | 0x0FFF0AA |
| R13 (SP) | 0x2000FD8 |
| R14 (LR) | 0xFFFFFFFF |
| R15 (PC) | 0x0000ACC |
| xPSR | 0x21000010 |
| Banked | |
| MSP | 0x2000FD8 |
| PSP | 0xBEF5DFF0 |
| System | |
| PRIMASK | 0 |
| CONTROL | 0x00 |
| Internal | |
| Mode | Handler |
| Stack | MSP |

Exception number 0x10
(interrupt number + 0x10)

EXC_RETURN

Mode changed to Handler.
Was already using MSP

4. Load PC With Addr of Exception Handler

| Memory Address | Value |
|-----------------|-----------------------|
| 0x0000_0000 | Initial Stack Pointer |
| 0x0000_0004 | Reset |
| 0x0000_0008 | NMI_IRQHandler |
| ... | |
| | IRQ0_Handler |
| | IRQ1_Handler |
| ... | |
| Reset: | code |
| ... | |
| NMI_IRQHandler: | code |
| ... | |
| IRQ0_Handler: | code |
| ... | |
| IRQ1_Handler: | code |

- The program counter is selected from the vector table depending on exception

Vector Table

| Exception number | IRQ number | Vector | Offset |
|--------------------|------------|------------|---------|
| 0 (initialization) | | Initial SP | 0x00 |
| 1 | | Reset | 0x04 |
| 2 | -14 | NMI | 0x08 |
| 3 | -13 | HardFault | 0x0C |
| 4 | | Reserved | 0x10 |
| 5 | | | |
| 6 | | | |
| 7 | | | |
| 8 | | | |
| 9 | | | |
| 10 | | | |
| 11 | -5 | SVCall | 0x2C |
| 12 | | Reserved | |
| 13 | | | |
| 14 | -2 | PendSV | 0x38 |
| 15 | -1 | SysTick | 0x3C |
| 16 | 0 | IRQ0 | 0x40 |
| 17 | 1 | IRQ1 | 0x44 |
| 18 | 2 | IRQ2 | 0x48 |
| . | | . | |
| 16+n | n | IRQn | 0x40+4n |

- Note:
LSB of the vector address is set to '1' to indicate that the handler uses Thumb code.

internal exceptions

external exceptions



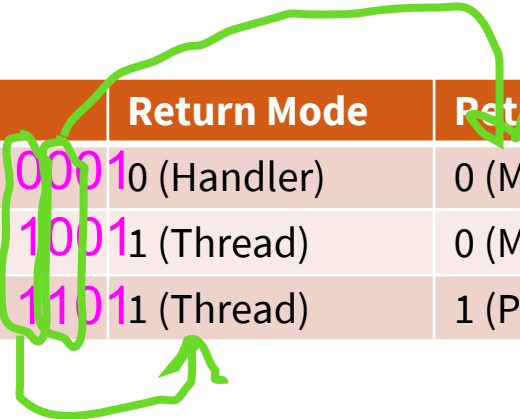
Upon Entry to Handler: uVision

| Registers | |
|-----------|------------|
| Register | Value |
| Core | |
| R0 | 0x00000000 |
| R1 | 0x00000000 |
| R2 | 0x00000000 |
| R3 | 0x0000008C |
| R4 | 0x40040000 |
| R5 | 0x00000ECC |
| R6 | 0x2EE00000 |
| R7 | 0x00000EAB |
| R8 | 0x9C7FB76F |
| R9 | 0x20000768 |
| R10 | 0x00000ECC |
| R11 | 0x00000ECC |
| R12 | 0x0FFFF0AA |
| R13 (SP) | 0x2000FD8 |
| R14 (LR) | 0xFFFFFFFF |
| R15 (PC) | 0x0000ACC |
| xPSR | 0x21000010 |
| Banked | |
| System | |
| Internal | |
| Mode | Handler |
| Stack | MSP |

```
42: void switch_isr(void) {  
0x00000ACC B510      PUSH      {r4,lr}
```

PC has been loaded
with start address of
handler

5. Load LR With EXC_RETURN Code



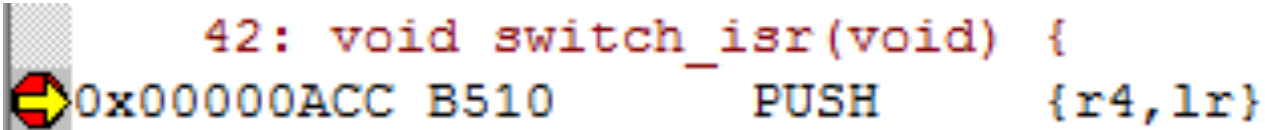
| EXC_RETURN | Return Mode | Return Stack | Description |
|-------------|-----------------|--------------|-----------------------------|
| 0xFFFF_FFF1 | 00010 (Handler) | 0 (MSP) | Return to exception handler |
| 0xFFFF_FFF9 | 10011 (Thread) | 0 (MSP) | Return to thread with MSP |
| 0xFFFF_FFFD | 11011 (Thread) | 1 (PSP) | Return to thread with PSP |

- EXC_RETURN value generated by CPU to provide information on how to return
 - Which SP to restore registers from? MSP (0) or PSP (1)
 - Previous value of SPSEL

most of time, MSP
 - Which mode to return to? Handler (0) or Thread (1)
 - Another exception handler may have been running when this exception was requested

6. Start Executing Exception Handler

- Exception handler starts running, unless preempted by a higher-priority exception
- Exception handler may save additional registers on stack
e.g. if handler may call a subroutine, LR and R4 must be saved



```
42: void switch_isr(void) {  
0x00000ACC B510      PUSH      {r4,lr}
```

Handler Has Saved More Context: uVision

42: void switch_isr(void) {
0x00000ACC B510 PUSH {r4,lr}

EXC_RETURN loaded to LR

Saved R4
Saved LR
Saved R0
Saved R1
Saved R2
Saved R3
Saved R12
Saved LR
Saved PC
Saved xPSR


SP value reduced since registers have been pushed onto stack

| Register | Value |
|----------|------------|
| R0 | 0x00000000 |
| R1 | 0x00000000 |
| R2 | 0x00000000 |
| R3 | 0x0000008C |
| R4 | 0x40040000 |
| R5 | 0x00000ECC |
| R6 | 0x2EE00000 |
| R7 | 0x00000EAB |
| R8 | 0x9C7F876F |
| R9 | 0x20000768 |
| R10 | 0x00000ECC |
| R11 | 0x00000ECC |
| R12 | 0x0FFFF0AA |
| R13 (SP) | 0x2000FD0 |
| R14 (LR) | 0xFFFFF9 |
| R15 (PC) | 0x0000ACE |
| xPSR | 0x21000010 |
| MSP | 0x2000FD0 |
| PSP | 0xBEF5DFF0 |
| PRIMASK | 0 |
| CONTROL | 0x00 |
| Mode | Handler |
| Stack | MSP |

| Address | Value |
|-----------|--|
| 0x2000FD0 | 40040000 FFFFFFF9 00000000 00000000 00000000 |
| 0x2000FE4 | 0000008C 0FFFFFF0AA 00000A0B 00000A70 21000000 |
| 0x2000FF8 | 00000002 00000A0B 000001F1 000001F1 000001F1 |
| 0x200100C | 000001F1 000001F1 000001F1 000001F1 000001F1 |
| 0x2001020 | 000001F1 000001F1 000001F1 000001F1 000001F1 |

Exiting an Exception Handler

Exiting an Exception Handler

1. Execute instruction triggering exception return processing in handler mode (EXC_RETURN)
2. Select return stack, restore context from that stack


The diagram shows two vertical green bars representing registers, labeled 'MSP' on the left and 'PSP' on the right. A green bracket is drawn above them, spanning from the top of the MSP bar to the top of the PSP bar, indicating a selection or restoration operation between the two stacks.
3. Resume execution of code at restored address

1. Execute Instruction for Exception Return

51: }
⇒ 0x00000B08 BD10 POP {r4, pc}

- No “return from interrupt” instruction
- Use regular instruction instead
 - **BX LR** - Branch to address in LR by loading PC with LR contents
 - **POP {..., PC}** - Pop address from stack into PC
- ... with a special value EXC_RETURN loaded into the PC to trigger exception handling processing
 - BX LR used if EXC_RETURN is still in LR
 - If EXC_RETURN has been saved on stack, then use POP

What Will Be Popped from Stack?

- R4: 0x4040_0000
- PC: 0xFFFF_FFF9

51: }
⇒ 0x00000B08 BD10 POP {r4,pc}

| Register | Value |
|----------|------------|
| Core | |
| R0 | 0x0000008C |
| R1 | 0x40040000 |
| R2 | 0xE000E280 |
| R3 | 0x0000008C |
| R4 | 0x40040000 |
| R5 | 0x00000ECC |
| R6 | 0x2EE00000 |
| R7 | 0x00000EAB |
| R8 | 0x9C7FB76F |
| R9 | 0x20000768 |
| R10 | 0x00000ECC |
| R11 | 0x00000ECC |
| R12 | 0x0FFFF0AA |
| R13 (SP) | 0x20000FD0 |
| R14 (LR) | 0x00000AE1 |
| R15 (PC) | 0x00000B08 |
| xPSR | 0x01000010 |
| Banked | |
| MSP | 0x20000FD0 |
| PSP | 0xBEF5DFF0 |
| System | |
| PRIMASK | 0 |
| CONTROL | 0x00 |
| Internal | |
| Mode | Handler |
| Stack | MSP |

| Memory 2 | |
|-------------|--|
| Address: | sp |
| 0x20000FD0: | 40040000 FFFFFFF9 00000000 00000000 00000000 |
| 0x20000FE4: | 0000008C 0FFFF0AA 00000A0B 00000A70 21000000 |
| 0x20000FF8: | 00000002 00000A0B 000001F1 000001F1 000001F1 |
| 0x2000100C: | 000001F1 000001F1 000001F1 000001F1 000001F1 |
| 0x20001020: | 000001F1 000001F1 000001F1 000001F1 000001F1 |

Saved R4

Saved LR

Saved R0

Saved R1

Saved R2

Saved R3

Saved R12

Saved LR

Saved PC

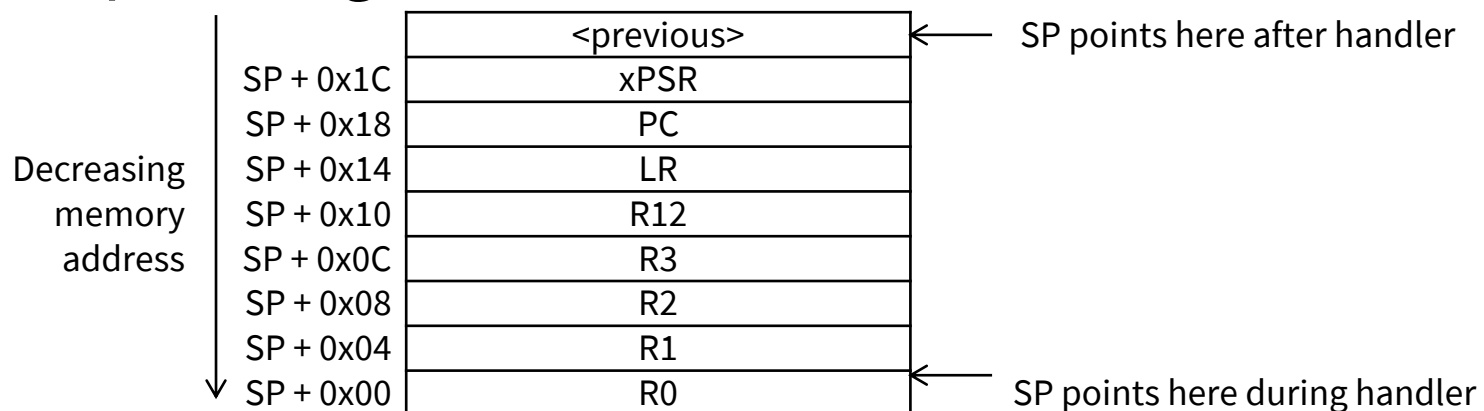
Saved xPSR

2. Select Stack, Restore Context

- Check EXC_RETURN (bit 2) to determine from which SP to pop the context

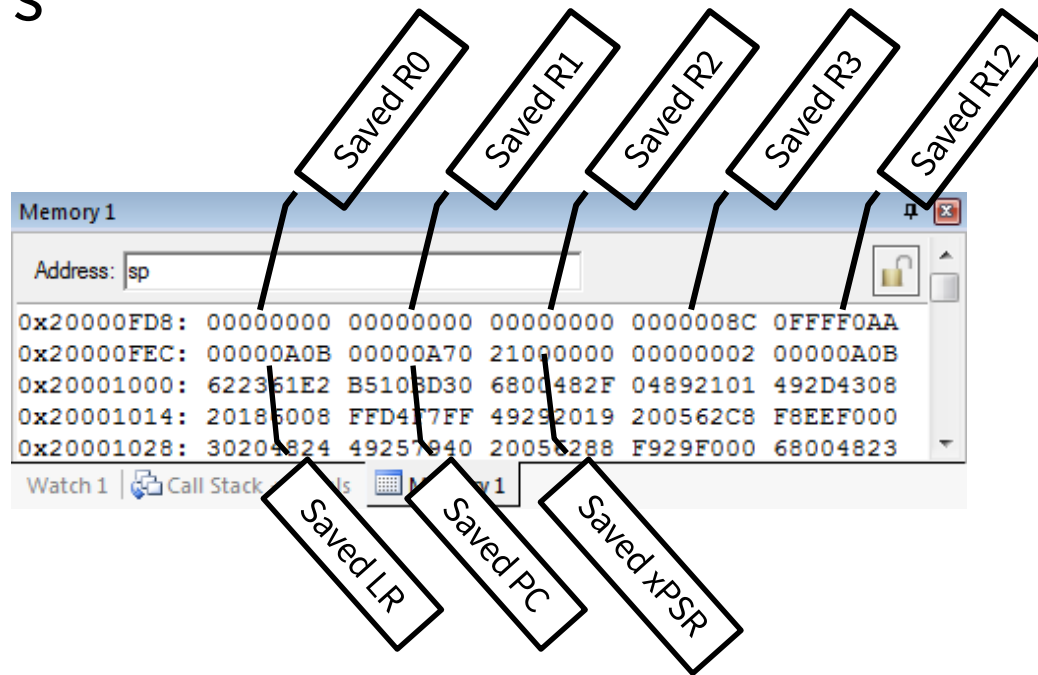
| EXC_RETURN | Return Stack | Description |
|-------------|--------------|--------------------------------------|
| 0xFFFF_FFF1 | 0 (MSP) | Return to exception handler with MSP |
| 0xFFFF_FFF9 | 0 (MSP) | Return to thread with MSP |
| 0xFFFF_FFFD | 1 (PSP) | Return to thread with PSP |

- Pop the registers from that stack



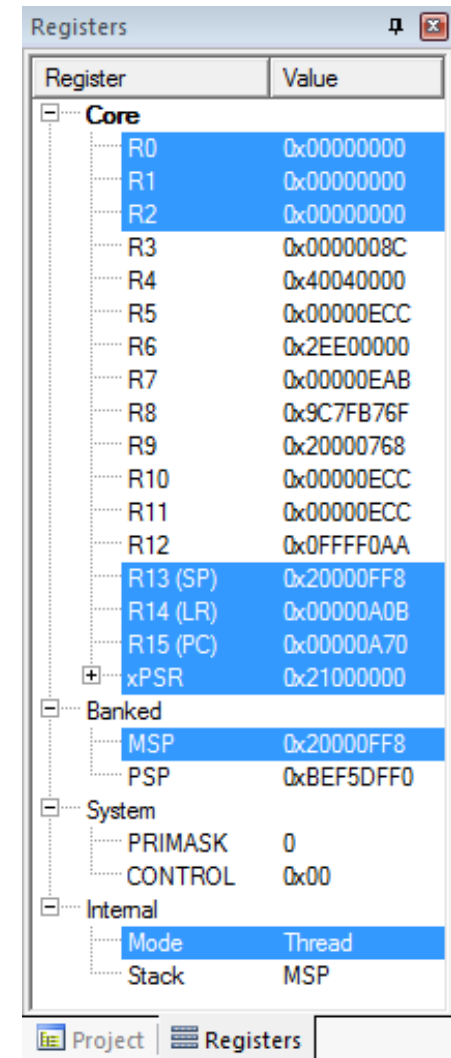
Example

- PC=0xFFFF_FFF9, so return to thread mode with main stack pointer
- Pop exception stack frame from stack back into registers



Resume Executing Previous Main Thread Code

- Exception handling registers have been restored: R0, R1, R2, R3, R12, LR, PC, xPSR
- SP is back to previous value
- Back in thread mode
- Next instruction to execute is at 0x0000_0A70



| Register | Value |
|-----------------|------------|
| Core | |
| R0 | 0x00000000 |
| R1 | 0x00000000 |
| R2 | 0x00000000 |
| R3 | 0x0000008C |
| R4 | 0x40040000 |
| R5 | 0x0000ECC |
| R6 | 0x2EE00000 |
| R7 | 0x0000EAB |
| R8 | 0x9C7FB76F |
| R9 | 0x20000768 |
| R10 | 0x0000ECC |
| R11 | 0x0000ECC |
| R12 | 0xFFFF0AA |
| R13 (SP) | 0x2000FF8 |
| R14 (LR) | 0x0000A0B |
| R15 (PC) | 0x0000A70 |
| xPSR | 0x21000000 |
| Banked | |
| MSP | 0x2000FF8 |
| PSP | 0xBEF5DFF0 |
| System | |
| PRIMASK | 0 |
| CONTROL | 0x00 |
| Internal | |
| Mode | Thread |
| Stack | MSP |

Processor Core Interrupts


Microcontroller Interrupts

- Hardware interrupts
 - Asynchronous: not related to what code the processor is currently executing
 - Examples: interrupt is asserted, character is received on serial port, or ADC converter finishes conversion
- Exceptions, Faults, software interrupts
 - Synchronous: are the result of specific instructions executing
 - Examples: undefined instructions, overflow occurs for a given instruction
 - We can enable and disable (mask) most interrupts as needed (maskable), others are non-maskable
- Interrupt service routine (ISR) **interrupt handler**
 - Subroutine which processor is forced to execute to respond to a specific event
 - After ISR completes, MCU goes back to previously executing code

inside NVIC, there are hundreds or thousands of interrupts. each interrupt is assigned a priority.

Nested Vectored Interrupt Controller

<https://developer.arm.com/documentation/ddi0337/h/nested-vectored-interrupt-controller?lang=en>



- NVIC manages and prioritizes external interrupts
- Interrupts are types of exceptions
 - Exceptions 16 through 16+N
- Modes
 - Thread Mode: entered on Reset
 - Handler Mode: entered on executing an exception
- Privilege level
- Stack pointers:
Main Stack Pointer (MSP) or Process Stack Pointer (PSP)
- Exception states: Inactive, Pending, Active, A&P

假如有个低一点优先级的 interrupt is triggered (这个时候没有更高优先级的 interrupt 了),突然有了一个高优先级的 interrupt。那么应该 abandon 低优先级, execute 高优先级。等到完成高优先级, resume 低优先级。

NVIC Registers and State

- Enable - Allows interrupt to be recognized by MCU
 - Accessed through two registers (set bits for interrupts)
Set enable with **NVIC_ISER**, clear enable with **NVIC_ICER**
 - CMSIS Interface: **NVIC_EnableIRQ(IRQnum)**,
NVIC_DisableIRQ(IRQnum)
- Pending - Interrupt has been requested but is not yet serviced
 - CMSIS:
NVIC_SetPendingIRQ(IRQnum),
NVIC_ClearPendingIRQ(IRQnum)

<https://interrupt.memfault.com/blog/arm-cortex-m-exceptions-and-nvic>

interruptmask: A central processing unit (CPU) feature that allows the computer to ignore (mask) an interrupt request until the mask bit is disabled.

Core Exception Mask Register

- Similar to “Global interrupt disable” bit in other MCUs
- PRIMASK - Exception mask register (CPU core)
 - Bit 0: PM Flag **PM: PRIMASK register** 除了non-maskable exceptions
 - Set to 1 to prevent activation of all exceptions with configurable priority
 - Clear to 0 to allow activation of all exception
 - Access using CPS, MSR and MRS instructions
 - Use to prevent data race conditions with code needing atomicity
- CMSIS-CORE API
 - **void __enable_irq()** - clears PM flag
 - **void __disable_irq()** - sets PM flag
 - **uint32_t __get_PRIMASK()** - returns value of PRIMASK
 - **void __set_PRIMASK(uint32_t x)** - sets PRIMASK to x

Prioritisation

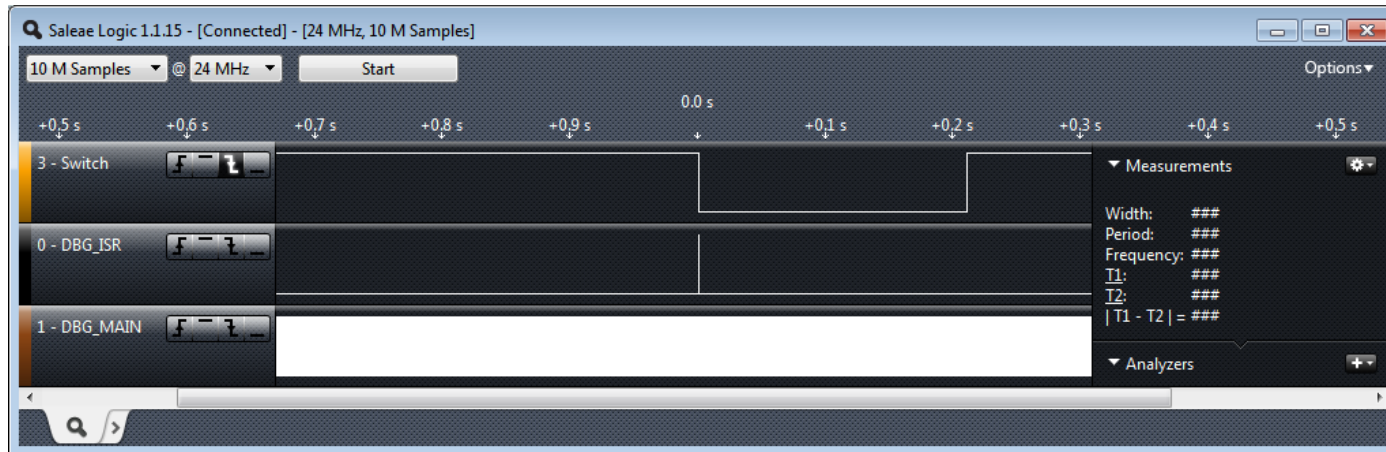
- Exceptions are prioritized to order the response simultaneous requests (smaller number = higher priority)
- Priorities of some exceptions are **fixed**
 - Reset: -3, highest priority
 - NMI: -2
 - Hard Fault: -1
- Priorities of other (peripheral) exceptions are **adjustable**
 - Value is stored in the interrupt priority register (IPR0-7)
 - 0x00
 - 0x40
 - 0x80
 - 0xC0

Special Cases of Prioritisation

- Simultaneous exception requests?
 - Lowest exception type number is serviced first
- New exception requested while a handler is executing?
 - New priority higher than current priority?
 - New exception handler **preempts** current exception handler
 - New priority lower than or equal to current priority?
 - New exception held in pending state
 - Current handler continues and completes execution
 - Previous priority level restored
 - New exception handled if priority level allows

Timing Analysis

Big Picture Timing Behavior



- Switch was pressed for about 0.21 s
 - ISR runs in response to switch signal's falling edge
 - Main seems to be running continuously (signal toggles between 1 and 0)
- Does it really? You will investigate this in the lab exercise.

Interrupt Response Latency

- Latency = time delay
- Why do we care?
 - This is overhead which wastes time, and increases as the interrupt rate rises
 - This delays our response to external events, which may or may not be acceptable for the application, such as sampling an analog waveform
- How long does it take?
 - Finish executing the current instruction or abandon it
 - Push various registers on to the stack, fetch vector
 - $C_{\text{IntResponseOvhd}}$: Overhead for responding to each interrupt
 - If we have external memory with wait states, this takes longer

Maximum Interrupt Rate

- We can only handle so many interrupts per second
 - $F_{\text{Max_Int}}$: maximum interrupt frequency
 - F_{CPU} : CPU clock frequency
 - C_{ISR} : Number of cycles ISR takes to execute
 - C_{Overhead} : Number of cycles of overhead for saving state, vectoring, restoring state, etc.
 - $F_{\text{Max_Int}} = F_{\text{CPU}} / (C_{\text{ISR}} + C_{\text{Overhead}})$
 - Note that model applies only when there is one interrupt in the system

Maximum Interrupt Rate

- When processor is responding to interrupts, it isn't executing our other code
 - U_{Int} : Utilization (fraction of processor time) consumed by interrupt processing
 - $U_{Int} = 100\% * F_{Int} * (C_{ISR} + C_{Overhead}) / F_{CPU}$
 - CPU looks like it's running the other code with CPU clock speed of $(1 - U_{Int}) * F_{CPU}$

$$U_{int} = 100\% * \frac{F_{int}}{F_{Max_int}}$$

CPU用全部精力处理interrupt的时候, $U_{int} = 100\%$

但是CPU还会用一些时间run the other code, 这个时候, interrupt的时间占有率就不是100%了。

Program Design with Interrupts

- How much work to do in ISR?
- Should ISRs re-enable interrupts?
- How to communicate between ISR and other threads?
 - Data buffering
 - Data integrity and race conditions

How much work to do in ISR?

- Trade-off: Faster response for ISR code will delay completion of other code
- In system with multiple ISRs with short deadlines, perform critical work in ISR and buffer partial results for later processing

ISR会花费额外的时间

Sharing Data Safely between ISRs and other Threads

Overview

- Volatile data – can be updated outside of the program's immediate control
- Non-atomic shared data – can be interrupted partway through read or write, is vulnerable to race conditions
`i++;` // is not thread safe because of multi-task concurrency.

Volatile Data

- Compilers assume that variables in memory do not change spontaneously, and optimize based on that belief
 - *Don't reload a variable from memory if current function hasn't changed it*
 - Read variable from memory into register (faster access)
 - Write back to memory at end of the procedure, or before a procedure call, or when compiler runs out of free registers
- This optimization can fail
 - Example: reading from input port, polling for key press
 - while (SW_0) ; will read from SW_0 once and reuse that value
 - Will generate an infinite loop triggered by SW_0 being true

这里的volatile和断电丢失的volatile不同。这里的volatile (易受损)是指容易被别的地方修改。

A variable should be declared **volatile** whenever its value could change unexpectedly.

In practice, only three types of variables could change:

1. Memory-mapped peripheral registers
2. Global variables modified by an interrupt service routine
3. Global variables within a multi-threaded application

The Volatile Directive

- Need to tell compiler which variables may change outside of its control
 - Use volatile keyword to force compiler to reload these vars from memory for each use *not cache or temporal memory*
volatile unsigned int num_ints;
 - Pointer to a volatile int
volatile int * var; // or
int volatile * var;
 - Now each C source read of a variable (e.g. status register) will result in an assembly language LDR instruction
 - Good explanation in Nigel Jones' "Volatile," *Embedded Systems Programming* July 2001

Non-Atomic Shared Data

- Want to keep track of current time and date
- Use 1 Hz interrupt from timer
- System
 - `current_time` structure tracks time and days since some reference event
 - `current_time`'s fields are updated by periodic 1 Hz timer ISR

```
void GetDateTime(DateTimeType * DT){  
    DT->day = current_time.day;  
    DT->hour = current_time.hour;  
    DT->minute = current_time.minute;  
    DT->second = current_time.second;  
}
```

```
void DateTimeISR(void){  
    current_time.second++;  
    if (current_time.second > 59){  
        current_time.second = 0;  
        current_time.minute++;  
        if (current_time.minute > 59) {  
            current_time.minute = 0;  
            current_time.hour++;  
            if (current_time.hour > 23) {  
                current_time.hour = 0;  
                current_time.day++;  
                ... //etc.  
            }  
        }  
    }  
}
```

Example: Checking the Time

- Problem
 - An interrupt at the wrong time will lead to half-updated data in DT
- Failure Case
 - current_time is {10, 23, 59, 59} (10th day, 23:59:59)
 - Task code calls GetDateTime(), which starts copying the current_time fields to DT: day = 10, hour = 23
 - A timer interrupt occurs, which updates current_time to {11, 0, 0, 0}
 - GetDateTime() resumes executing, copying the remaining current_time fields to DT: minute = 0, second = 0
 - DT now has a time stamp of {10, 23, 0, 0}.
 - The system thinks time just jumped backwards one hour!
- Fundamental problem – “race condition”
 - Preemption enables ISR to interrupt other code and possibly overwrite data
 - Must ensure atomic (indivisible) access to the object
 - Native atomic object size depends on processor’s instruction set and word size.
 - Is 32 bits for ARM

看看这个例子

Examining the Problem More Closely

- Must protect any data object which both:
 - Requires multiple instructions to read or write (non-atomic access), and
 - Is potentially written by an ISR
- How many tasks/ISRs can write to the data object?
 - One? Then we have one-way communication
 - Must *ensure the data isn't overwritten* partway through being *read*
 - Writer and reader don't interrupt each other
 - More than one?
 - Must *ensure the data isn't overwritten* partway through being *read*
 - Writer and reader don't interrupt each other
 - Must *ensure the data isn't overwritten* partway through being *written*
 - Writers don't interrupt each other

Definitions

- **Race condition:** Anomalous behavior due to unexpected critical dependence on the relative timing of events. Result of example code depends on the relative timing of the read and write operations.
- **Critical section:** A section of code which creates a possible race condition. The code section can only be executed by one process at a time. Some synchronization mechanism is required at the entry and exit of the critical section to ensure exclusive use.

Solution: Briefly Disable Preemption

- Prevent preemption within critical section
- If an ISR can *write* to the shared data object, need to *disable interrupts*
 - save current interrupt masking state in m
 - disable interrupts
- Restore *previous state* afterwards (interrupts may have already been disabled for another reason)
- Use CMSIS-CORE to save, control and restore interrupt masking state
- Avoid if possible
 - Disabling interrupts delays response to all other processing requests
 - Make this time as short as possible (e.g. a few instructions)

```
void GetDateTime(DateTimeType * DT){
    uint32_t m;

    m = __get_PRIMASK();
    __disable_irq();

    DT->day = current_time.day;
    DT->hour = current_time.hour;
    DT->minute = current_time.minute;
    DT->second = current_time.second;
    __set_PRIMASK(m);
}
```

Summary for Sharing Data

- In thread/ISR diagram, identify shared data
- Determine which shared data is too large to be handled atomically by default
This needs to be protected from preemption (e.g. disable interrupt(s), use an RTOS synchronization mechanism)
- Declare (and initialize) shared variables as volatile in main file (or globals.c)
 - **volatile int my_shared_var=0;** don't go to RAM or other temporary memory
- Update extern.h to make these variables available to functions in other files
 - **volatile int my_shared_var;**
 - **#include "extern.h"** in every file which uses these shared variables
- When using long (non-atomic) shared data, save, disable and restore interrupt masking status
 - CMSIS-CORE interface: **__disable_irq(), __get_PRIMASK(), __set_PRIMASK()**